

Boosting is an ensemble learning technique that focuses on **reducing bias** and improving model accuracy. It builds models **sequentially**, where each new model focuses on correcting the errors made by the previous models. The final predictions are combined in a weighted manner, giving more importance to the stronger models.

✓ Steps of Boosting

1. Train Initial Model: Train the first weak learner (e.g., a decision tree) on the dataset.
2. Identify Errors: Evaluate the errors (misclassified samples) of the current model.
3. Update Weights: Increase the importance (weights) of the misclassified samples so that the next model focuses more on them.
4. Combine Models: Sequentially add new models, and their predictions are combined based on their performance.

Why Boosting Works?

1. Bias Reduction: Sequentially correcting errors minimizes bias.
2. Weighted Contributions: Stronger models have more influence, improving final predictions.
3. Flexibility: Boosting adapts well to various data distributions.

✓ Common Boosting Algorithms

1. AdaBoost: Adjusts weights of misclassified samples.
2. Gradient Boosting: Minimizes the loss function by training models on the residual errors.
3. XGBoost/LightGBM: Optimized versions of gradient boosting.

A weak learner is a simple model that performs slightly better than random guessing. Here, we use a Decision Tree with a maximum depth of 1. This means:

- The tree can only make one split, resulting in a very simple decision boundary.
- This is why it's called a "shallow" decision tree.

Example of Boosting in Action

Imagine the dataset has 3 misclassified samples (A, B, and C):

1. First Learner: Correctly classifies most samples but misclassifies A and B.
2. Second Learner: Focuses on A and B, getting them right but misclassifies C.
3. Third Learner: Focuses on C and corrects it. Finally, the ensemble combines all predictions, ensuring that errors from earlier models are corrected by later models.

✓ Training and Prediction

Boosting happens sequentially, so the model works in the following way:

1. Train the First Weak Learner:
 - Train the decision tree on the dataset.
 - Evaluate the errors (misclassified points).
2. Focus on Errors:
 - Increase the weight of misclassified samples so they are given more importance.
 - This ensures the next weak learner focuses more on these difficult cases.
3. Repeat for All Learners:
 - Each new weak learner is trained on the updated dataset (with new weights).
 - This process continues for all 50 weak learners.
4. Combine Results:

- Each learner's prediction is weighted based on its accuracy.
- The final prediction is a weighted combination of all weak learners.

We are using the Iris dataset for a classification example with AdaBoostClassifier.

```
# Import required libraries
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.ensemble import AdaBoostClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score
```

```
# Load the Iris dataset
```

```
iris = load_iris()
X, y = iris.data, iris.target
```

```
# Split data into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
```

```
# Initialize a base model (Decision Tree with max depth)
base_model = DecisionTreeClassifier(max_depth=1, random_state=42)
```

```
# Create an AdaBoostClassifier
```

```
adaboost_model = AdaBoostClassifier(
    estimator=base_model,      # Base learner
    n_estimators=50,           # Number of weak learners
    learning_rate=1.0,         # Controls the contribution of each learner
    random_state=42
)
```

▼ AdaBoostClassifier Parameters

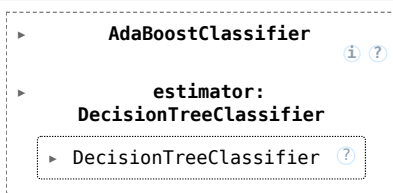
1. n_estimators=50

- We train 50 weak learners (decision trees in this case) sequentially.
- Each weak learner focuses on fixing the errors made by the previous ones.
- Impact: Increasing n_estimators can lead to better performance but may risk overfitting if set too high.

2. learning_rate=1.0

- This controls how much influence each weak learner has on the final prediction.
- A lower learning rate means the model learns more gradually (but may require more estimators).
- Impact: It adjusts the weight updates during boosting. A higher learning rate increases the contribution of each learner but risks instability.

```
# Train the AdaBoost model
adaboost_model.fit(X_train, y_train)
```



```
# Make predictions
y_pred = adaboost_model.predict(X_test)
```

```
# Evaluate accuracy
```

```
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy of AdaBoost Classifier: {accuracy:.2f}")
```

```
↗ Accuracy of AdaBoost Classifier: 1.00
```

Code Walkthrough

1. Base Model: A shallow decision tree (max depth=1) is used as a weak learner.
2. AdaBoostClassifier Parameters:
 - n_estimators=50: Use 50 weak learners sequentially.
 - learning_rate=1.0: Controls how much each model contributes to the final prediction.
3. Training and Prediction: Sequentially trains models, focusing on correcting previous errors.

Here's an example of using Gradient Boosting for classification with the Iris dataset. We are using GradientBoostingClassifier from Scikit-learn.

▽ Gradient Boosting

Gradient Boosting is a boosting technique where each weak learner is trained to predict the **residual errors (differences between the true values and predictions)** of the previous learners. The model minimizes a loss function (e.g., log-loss for classification or mean squared error for regression) using gradient descent.

```
# Import required libraries
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.metrics import accuracy_score
```

```
# Load the Iris dataset
iris = load_iris()
X, y = iris.data, iris.target
```

```
# Split data into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
```

```
# Create a GradientBoostingClassifier
gb_model = GradientBoostingClassifier(
    n_estimators=100,      # Number of weak learners
    learning_rate=0.1,     # Step size for weight updates
    max_depth=3,          # Maximum depth of each decision tree
    random_state=42
)
```

```
# Train the Gradient Boosting model
gb_model.fit(X_train, y_train)
```

```
↗ GradientBoostingClassifier ⓘ ?
GradientBoostingClassifier(random_state=42)
```

```
# Make predictions
y_pred = gb_model.predict(X_test)
```

```
# Evaluate accuracy
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy of Gradient Boosting Classifier: {accuracy:.2f}")
```

↻ Accuracy of Gradient Boosting Classifier: 1.00

✓ Code Walkthrough

1. Model Initialization:

- `n_estimators=100`: Use 100 decision trees in sequence.
- `learning_rate=0.1`: Each tree contributes less, allowing for gradual improvement.
- `max_depth=3`: Restrict trees to a maximum depth of 3 for simplicity and to avoid overfitting.

2. Training:

- The model trains sequentially, where each tree focuses on correcting the residual errors of the previous trees.

3. Prediction:

- Final predictions are made by combining the results of all trees.

4. Evaluation:

- The accuracy of the model is computed on the test set.

Unlike AdaBoost, where you specify a weak learner explicitly, Gradient Boosting uses shallow decision trees (typically one with a small maximum depth) internally to fit the residuals. You control these weak learners indirectly using parameters like `max_depth`, `n_estimators`, and `learning_rate`.

How Gradient Boosting Works Internally

1. Weak Learners:

- By default, the weak learners are shallow decision trees (max depth controlled by `max_depth`).
- These trees are trained sequentially.

2. Residual Learning:

- The first tree is trained on the original target values.
- Each subsequent tree is trained on the residual errors (differences between true and predicted values from all previous trees).

3. Weighted Predictions:

- The predictions from all trees are combined (weighted by the learning rate) to make the final prediction.

Why Gradient Boosting Works

1. Error Correction: Each tree is trained on the residuals, incrementally reducing the error.
2. Flexibility: Can optimize for various loss functions (e.g., classification, regression).
3. Hyperparameters: Fine-tuning parameters like `n_estimators`, `learning_rate`, and `max_depth` helps control performance and avoid overfitting.